

# Mining Twitter for Social Event and Misinformation Detection

**Yingtong Dou**

Department of Computer Science

University of Illinois at Chicago

Homepage: <http://ytongdou.com>

Email: [ydou5@uic.edu](mailto:ydou5@uic.edu)

Twitter: [@dozee\\_sim](https://twitter.com/dozee_sim)



UNIVERSITY OF  
**ILLINOIS CHICAGO**

# Outlines

- **Background:** Twitter and Graph Neural Network
- **Paper I:** Detecting Misinformation on Twitter
- **Paper II:** Discovering Social Event on Twitter
- **Q&A**

# Twitter Fact Sheet

- Twitter has **396M** users in 2021.
- **51M** Americans use Twitter daily.
- In the US, **92%** of tweets come from the top **10%** of users.
- **48%** Twitter users use Twitter to get news.
- Every second, there are **6k** new tweets are tweeted on Twitter.



# Twitter vs FB, Reddit, TikTok

- Twitter has an explicit social network structure.
- Twitter is an open platform to discuss public events.
- Structured Twitter data can be accessible via Twitter Developer API.




Image from Jabeur, et al. "Uprising microblogs: A Bayesian network retrieval model for tweet search." *ACM symposium on applied computing*. 2012.

# Leveraging Twitter Data with Twitter API

- **Conduct academic research**
  - From computer science to social science.
- **Solve problems with applied research for NGOs**
  - Conduct scientific studies that solve problems to impact the mission of NGOs.
- **Enrich investigative journalism and independent research**
  - Use Twitter data to explore global to local topics and events that can inform projects and publications.
- **Conduct market research for business**
  - Understand your audience and what they value by uncovering trends and surfacing important conversations on Twitter.


# CS Research using Twitter Data

- Social Event Detection
- Fraudster/Spammer/Bot/Sybil Detection
- Social Recommendation
- Misinformation Detection
- Graph Mining (link prediction, node classification/clustering)
- Sentiment/Emotion/Opinion/Stance/Topic Mining
- Multi-modal Data Mining
- Conversational Agent
- And More ...

# Botnet on Twitter




Human-like bot accounts on Twitter.




The profile pictures of accounts on the left, made opaque and superimposed.

# Botnet on Twitter

- A group of Twitter accounts retweet each other's tweets with similar content
- The densely connected accounts can be easily discovered from the graph perspective.



# Graph Mining and Graph Neural Network



**Key idea:** the connected nodes are similar (homophily assumption).  
**Objective:** learn an optimal neural network-based neighborhood encoder.

# Graph Neural Network for Graph Classification




Image from Ying, Zhitao, et al. "Hierarchical graph representation learning with differentiable pooling." NeurIPS (2018).

# In This Talk

- **Misinformation Detection on Twitter**
  - Introduce how to crawl Twitter data.
  - Modeling news propagation on Twitter as a tree-structured graph.
  - Supervised graph classification problem.
- **Social event detection on twitter**
  - Encoding different relations and entities on Twitter using GNN.
  - Unsupervised node clustering problem.

# SIGIR'21: Misinformation Detection on Twitter

## User Preference-aware Fake News Detection

**Yingtong Dou, Congying Xia, Philip Yu (University of Illinois at Chicago)**

**Kai Shu (Illinois Institute of Technology)**

**Lichao Sun (Lehigh University)**

Paper: <https://arxiv.org/pdf/2104.12259.pdf>

Code: <https://github.com/safe-graph/GNN-FakeNews>

Benchmark: <https://paperswithcode.com/dataset/upfd>

PyG Example: <https://tinyurl.com/a6s92t37>

DGL Example: <https://tinyurl.com/yjwvd93b>



UNIVERSITY OF  
**ILLINOIS CHICAGO**

ILLINOIS INSTITUTE  
OF TECHNOLOGY




**LEHIGH**  
UNIVERSITY  
11

# Problem Description

- **Fake news detection**
  - Fake news and real news are circulating on the social network.
  - A fake news has suspicious signals from its content, source, and social context.
  - Given the annotated news and their metadata, we want to train a neural network to classify unlabeled news.
- **Our motivation**
  - A social network user has his/her preference in consuming news. For a given piece of news, we assume its engaged users prefer similar types/content of news.
  - We propose to model user news consumption preference using user historical Tweets.



# Data Collection

- **Existing Data:**
  - The [FakeNewsNet](#) dataset from Prof. Kai Shu.
  - News content, timestamp, news retweets id, retweets timestamp.
- **We need to collect:**
  - Retweet user metadata.
  - Retweet user historical tweets.



# Twitter Developer API

- University student can apply a free Twitter developer account using the school email.
- Twitter API can control a Twitter account using APIs and can access most of Twitter data in a limited request rate.



# Python Crawler Code

- We use a Python Twitter crawler called [tweepy](#) to crawl the Twitter data.
- Tweepy will wait when API meets the rate limit.
- `api.user_timeline` will return a [Twitter status object](#) in json format.

```
import tweepy
import json

# Twitter Developer API tokens provided by Twitter
auth = tweepy.OAuthHandler('xxx', 'xxx')
auth.set_access_token('xxx', 'xxx')

api = tweepy.API(auth, wait_on_rate_limit=True, wait_on_rate_limit_notify=True)

m, n = 0, 0
for i, user in enumerate(id_mappings): # user.id to twitter.id mappings
    try:
        # get recent 200 tweets of the user
        statuses = api.user_timeline(user_id=user, count=200)
        json_object = [json.dumps(s._json) + '\n' for s in statuses]
        # write the recent 200 tweets into a json file
        with open('content/' + str(user) + ".json", "w") as outfile:
            outfile.writelines(json_object)
        outfile.close()
    except tweepy.TweepError as err: # handled deleted/suspended accounts
        if str(err) == 'Not authorized.':
            m += 1
            print(f'Not authorized: {m}')
        else:
            n += 1
            print(f'Page does not exist: {n}')
    print(f'user number: {i}')

print(f'Not authorized: {m}, Page does not exist: {n}.')
```

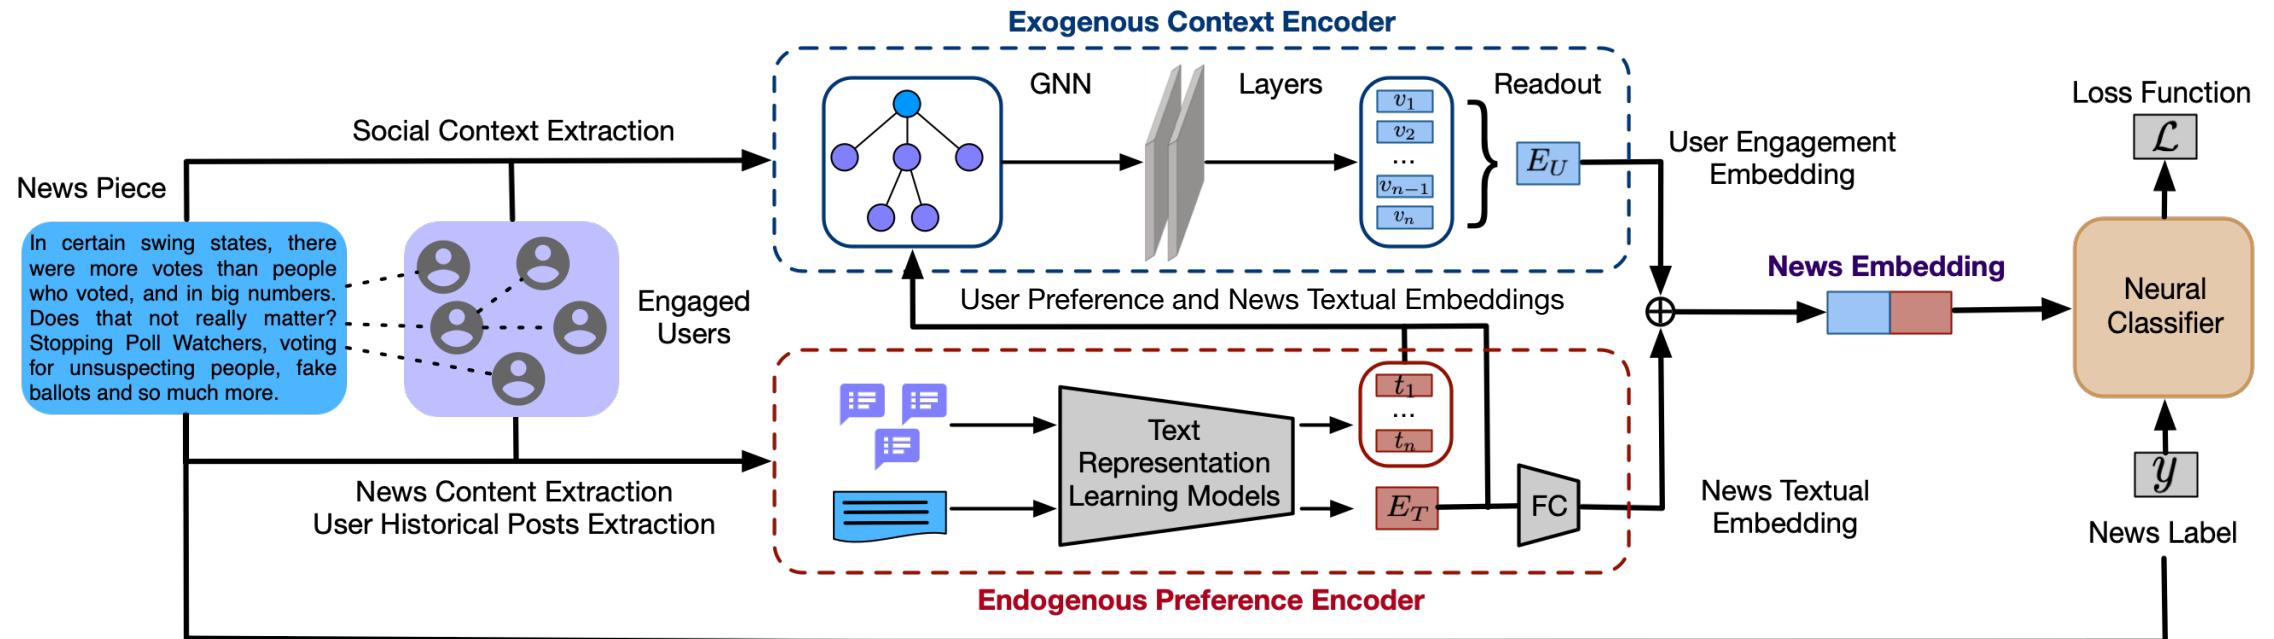
# Encoding User Preference

- **Preprocessing text data**
  - Remove special characters and emojis
  - Remove the “#” and “@”
  - Combining all tweets as a document
- **Three text encoders**
  - TFIDF embedding learned from current corpus.
  - Pretrained Word2vec vectors (using [spaCy](#)).
  - Pretrained BERT.

# Baseline: Profile-based Features

- 1) Verified?, 2) Enable geo-spatial positioning,
- 3) Followers count, 4) Friends count,
- 5) Status count, 6) Favorite count, 7) Number of lists,
- 8) Created time (No. of months since Twitter established),
- 9) Number of words in the description,
- 10) Number of words in the screen name

# Proposed UPFD Model



# Experiment Setup

Data	#Graphs	#Fake News	#Total Nodes	#Total Edges	#Avg. Nodes per Graph
Politifact	314	157	41,054	40,740	131
Gossipcop	5464	2732	314,262	308,798	58


- **Baselines**
  - SAFE
  - CSI
  - GNN-CL
  - GCNFN
- **Metrics**
  - Accuracy
  - F1-score
- **Implementation**
  - Python
  - [PyTorch](#)
  - [PyTorch\\_Geometric](#)

# Experiment Result

**Table 2: The fake news detection performance of baselines and our model. Stars denote statistically significant under the t-test ( $* p \leq 0.05$ ,  $** p \leq 0.01$ ,  $*** p \leq 0.001$ ).**

	<b>Model</b>	<b>POL</b>		<b>GOS</b>	
		ACC	F1	ACC	F1
News Only	SAFE [36]	73.30	72.87	77.37	77.19
	CSI [23]	76.02	75.99	75.20	75.01
	BERT+MLP	71.04	71.03	85.76	85.75
	word2vec+MLP	76.47	76.36	84.61	84.59
News + User	GNN-CL [8]	62.90	62.25	95.11	95.09
	GCNFN [17]	83.16	83.56	96.38	96.36
	UPFD (ours)	84.62*	84.65*	97.23**	97.22***

Proposed UPFD model has the best performance.



**Figure 2: The fake news detection performance of different variants of UPFD framework. -END/-EXO represents the UPFD variant without endogenous/exogenous information.**

Encoding user preference could improve fake news detection performance.

# Future Directions

- Encode fine-grained user preference signals.
- Explain fake news detection results based on user preference.
- Consider the temporal information.
- Fake news early detection.
- News popularity/engagement prediction.

# WWW'21: Social Event Detection on Twitter

## Knowledge-Preserving Incremental Social Event Detection via Heterogeneous GNNs

Yuwei Cao<sup>[1]</sup>, Hao Peng<sup>[2]</sup>, Jia Wu<sup>[3]</sup>, Yingtong Dou<sup>[1]</sup>, Jianxin Li<sup>[2]</sup>, Philip S. Yu<sup>[1]</sup>

Paper: <https://arxiv.org/pdf/2101.08747.pdf>

Code: <https://github.com/RingBDStack/KPGNN>



THE WEB  
CONFERENCE

[1]



[2]



[3]



# Problem Description

- Social event (e.g., the Notre-Dame Cathedral fire) reflect group social behaviors and wide-spread public concerns.
- Social event detection has many applications in fields including crisis management, product recommendation, and decision making.
- Social event detection can be formalized as **extracting clusters of co-related messages from social streams (i.e., sequences of social media messages)** to represent events.

# Challenges and Our Solution

- **Challenges**
  - We should leverage the rich semantic and relational information on Twitter.
  - The model should acquire, preserve, and extend knowledge given the streaming social messages.
- **Our solution:**
  - Use GNNs to learn the semantic and structural information.
  - Propose an incremental learning framework which has the pre-training, detection, and maintenance stages.

# Dataset Information

- **Twitter Event Corpus**
  - 68,841 manually labeled tweets related to 503 event classes, spreading over a period of four weeks.
  - 'event\_id': manually labeled event class
  - 'tweet\_id': tweet id 'text': content of the tweet
  - 'created\_at': timestamp of the tweet
  - 'user\_id': the id of the sender
  - 'user\_loc', 'place\_type', 'place\_full\_name': the location of the sender
  - 'hashtags': hashtags contained in the tweet
  - 'user\_mentions': user mentions contained in the tweet
  - 'image\_urls': links to the images contained in the tweet
  - 'entities': a list, named entities in the tweet (extracted using spaCy)
  - 'words': a list, tokens of the tweet (hashtags and user mentions are filtered out)
  - 'filtered\_words': a list, lower-cased words of the tweet
  - 'sampled\_words': a list, sampled words of the tweet

# Proposed Method - KPGNN

Unsupervised  
Loss Function

## Raw Messages


m1  
@user1 Apr 15th, 2019  
Still in tears when  
remembering the horrible  
fire @user2

m2  
@user2 Apr 15th, 2019  
The cultural and historic  
impact of Paris' Notre  
Dame Cathedral <https://...>



m3  
@user3 Apr 16th, 2019  
Restore the Notre Dame  
Cathedral after the fire  
[#notredamefire](https://...)

## Preprocessing

**(a) Heterogeneous Social Graph**




**(c) Homogeneous Message Graph**




**(b) Initial Message Features**

## Message Embedding



**(d) Message Embedding via GNN**


## Training



**(e) Contrastive Learning**

## Detection

social events



**(f) Message Clustering**

**Graph Construction**

**GNN Training**

**Detection Event  
By Clustering**


# KPGNN – Graph Construction

## Raw Messages


- m1  
@user1 Apr 15th, 2019  
Still in tears when  
remembering the horrible  
fire @user2
- m2  
@user2 Apr 15th, 2019  
The cultural and historic  
impact of Paris' Notre  
Dame Cathedral <https://...>
- m3  
@user3 Apr 16th, 2019  
Restore the Notre Dame  
Cathedral after the fire  
[#notredamefire](https://...)
- ...

## Preprocessing

### (a) Heterogeneous Social Graph




### (c) Homogeneous Message Graph



### (b) Initial Message Features

- (a): word nodes; name entity nodes; user nodes; tweet message nodes.
- (a) → (c): If two message have **at least one common neighbor**, we add one edge between them.
- (b) Message feature vector contains natural language semantics and temporal information.

# KPGNN -- Incremental Learning Lifecycle



## Nodes Removing Strategy

1. Keep all nodes.
2. Keep the relevant nodes.
3. Keep the latest nodes.

# Experiment Setup

- **Baselines**
  - Word2vec
  - LDA
  - WMD
  - BERT
  - Bi-LSTM
  - PP-GCN
  - EventX
- **Metrics**
  - NMI
  - AMI
  - ARI
- **Implementation**
  - Python
  - PyTorch
  - Deep Graph Library

# Experiment Result

**Table 2: Offline evaluation results on the Twitter dataset.** The best results are marked in bold and second-best in italic.

Metrics	Word2vec [26]	LDA [3]	WMD [20]	BERT [6]	BiLSTM [12]	PP-GCN [28]	EventX [21]	KPGNN <sub>t</sub>	KPGNN
NMI	.44±.00	.29±.00	.65±.00	.64±.00	.63±.00	.68±.02	<b>.72±.00</b>	.69±.01	.70±.01
AMI	.13±.00	.04±.00	.50±.00	.44±.00	.41±.00	.50±.02	.19±.00	.51±.00	<b>.52±.01</b>
ARI	.02±.00	.01±.00	.06±.00	.07±.00	.17±.00	.20±.01	.05±.00	.21±.01	<b>.22±.01</b>






## Performance on static dataset

**Table 4: The statistics of the social stream.**

Blocks	$M_0$	$M_1$	$M_2$	$M_3$	$M_4$	$M_5$	$M_6$	$M_7$	$M_8$	$M_9$	$M_{10}$
# of messages	20,254	8,722	1,491	1,835	2,010	1,834	1,276	5,278	1,560	1,363	1,096
Blocks	$M_{11}$	$M_{12}$	$M_{13}$	$M_{14}$	$M_{15}$	$M_{16}$	$M_{17}$	$M_{18}$	$M_{19}$	$M_{20}$	$M_{21}$
# of messages	1,232	3,237	1,972	2,956	2,549	910	2,676	1,887	1,399	893	2,410

## Incremental evaluation setting

# Different Nodes Removing Strategies



- All messages
- Relevant messages
- ▲— Latest messages

# Thank You!

## Q & A

**Yingtong Dou**

Department of Computer Science

University of Illinois at Chicago

Homepage: <http://ytongdou.com>

Email: [ydou5@uic.edu](mailto:ydou5@uic.edu)

Twitter: [@dozee\\_sim](https://twitter.com/dozee_sim)

